

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df = pd.read_csv("cities.csv")
df
```

```
Out[2]:
```

| | id | name | state_id | state_code | state_name | country_id | country_code | country_name |
|--------|--------|---------------------|----------|------------|-------------------|------------|--------------|--------------|
| 0 | 52 | Ashkāsham | 3901 | BDS | Badakhshan | 1 | AF | Afghanistan |
| 1 | 68 | Fayzabad | 3901 | BDS | Badakhshan | 1 | AF | Afghanistan |
| 2 | 78 | Jurm | 3901 | BDS | Badakhshan | 1 | AF | Afghanistan |
| 3 | 84 | Khandūd | 3901 | BDS | Badakhshan | 1 | AF | Afghanistan |
| 4 | 115 | Rāghistān | 3901 | BDS | Badakhshan | 1 | AF | Afghanistan |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 150449 | 131496 | Redcliff | 1957 | MI | Midlands Province | 247 | ZW | Zimbabwe |
| 150450 | 131502 | Shangani | 1957 | MI | Midlands Province | 247 | ZW | Zimbabwe |
| 150451 | 131503 | Shurugwi | 1957 | MI | Midlands Province | 247 | ZW | Zimbabwe |
| 150452 | 131504 | Shurugwi District | 1957 | MI | Midlands Province | 247 | ZW | Zimbabwe |
| 150453 | 131508 | Zvishavane District | 1957 | MI | Midlands Province | 247 | ZW | Zimbabwe |

150454 rows × 11 columns



```
In [3]: df.head()
```

```
Out[3]:
```

| | id | name | state_id | state_code | state_name | country_id | country_code | country_name | latitude |
|---|-----|-----------|----------|------------|------------|------------|--------------|--------------|----------|
| 0 | 52 | Ashkāsham | 3901 | BDS | Badakhshan | 1 | AF | Afghanistan | 36.68333 |
| 1 | 68 | Fayzabad | 3901 | BDS | Badakhshan | 1 | AF | Afghanistan | 37.11664 |
| 2 | 78 | Jurm | 3901 | BDS | Badakhshan | 1 | AF | Afghanistan | 36.86477 |
| 3 | 84 | Khandūd | 3901 | BDS | Badakhshan | 1 | AF | Afghanistan | 36.95127 |
| 4 | 115 | Rāghistān | 3901 | BDS | Badakhshan | 1 | AF | Afghanistan | 37.66079 |



Data cleaning and pre processing

In [4]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150454 entries, 0 to 150453
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               150454 non-null  int64
1   name            150454 non-null  object
2   state_id        150454 non-null  int64
3   state_code      150129 non-null  object
4   state_name      150454 non-null  object
5   country_id      150454 non-null  int64
6   country_code    150406 non-null  object
7   country_name    150454 non-null  object
8   latitude        150454 non-null  float64
9   longitude       150454 non-null  float64
10  wikiDataId      147198 non-null  object
dtypes: float64(2), int64(3), object(6)
memory usage: 12.6+ MB
```

In [5]:

```
df.describe()
```

Out[5]:

| | id | state_id | country_id | latitude | longitude |
|--------------|---------------|---------------|---------------|---------------|---------------|
| count | 150454.000000 | 150454.000000 | 150454.000000 | 150454.000000 | 150454.000000 |
| mean | 76407.091689 | 2678.377677 | 140.658460 | 31.556175 | 2.369557 |
| std | 44357.755335 | 1363.513591 | 70.666123 | 22.813220 | 68.012770 |
| min | 1.000000 | 1.000000 | 1.000000 | -75.000000 | -179.121980 |
| 25% | 38160.250000 | 1451.000000 | 82.000000 | 19.000000 | -58.468150 |
| 50% | 75975.500000 | 2174.000000 | 142.000000 | 40.684720 | 8.669980 |
| 75% | 115204.750000 | 3905.000000 | 207.000000 | 47.239220 | 27.750000 |
| max | 153528.000000 | 5116.000000 | 247.000000 | 73.508190 | 179.466000 |

In [6]:

```
df.columns
```

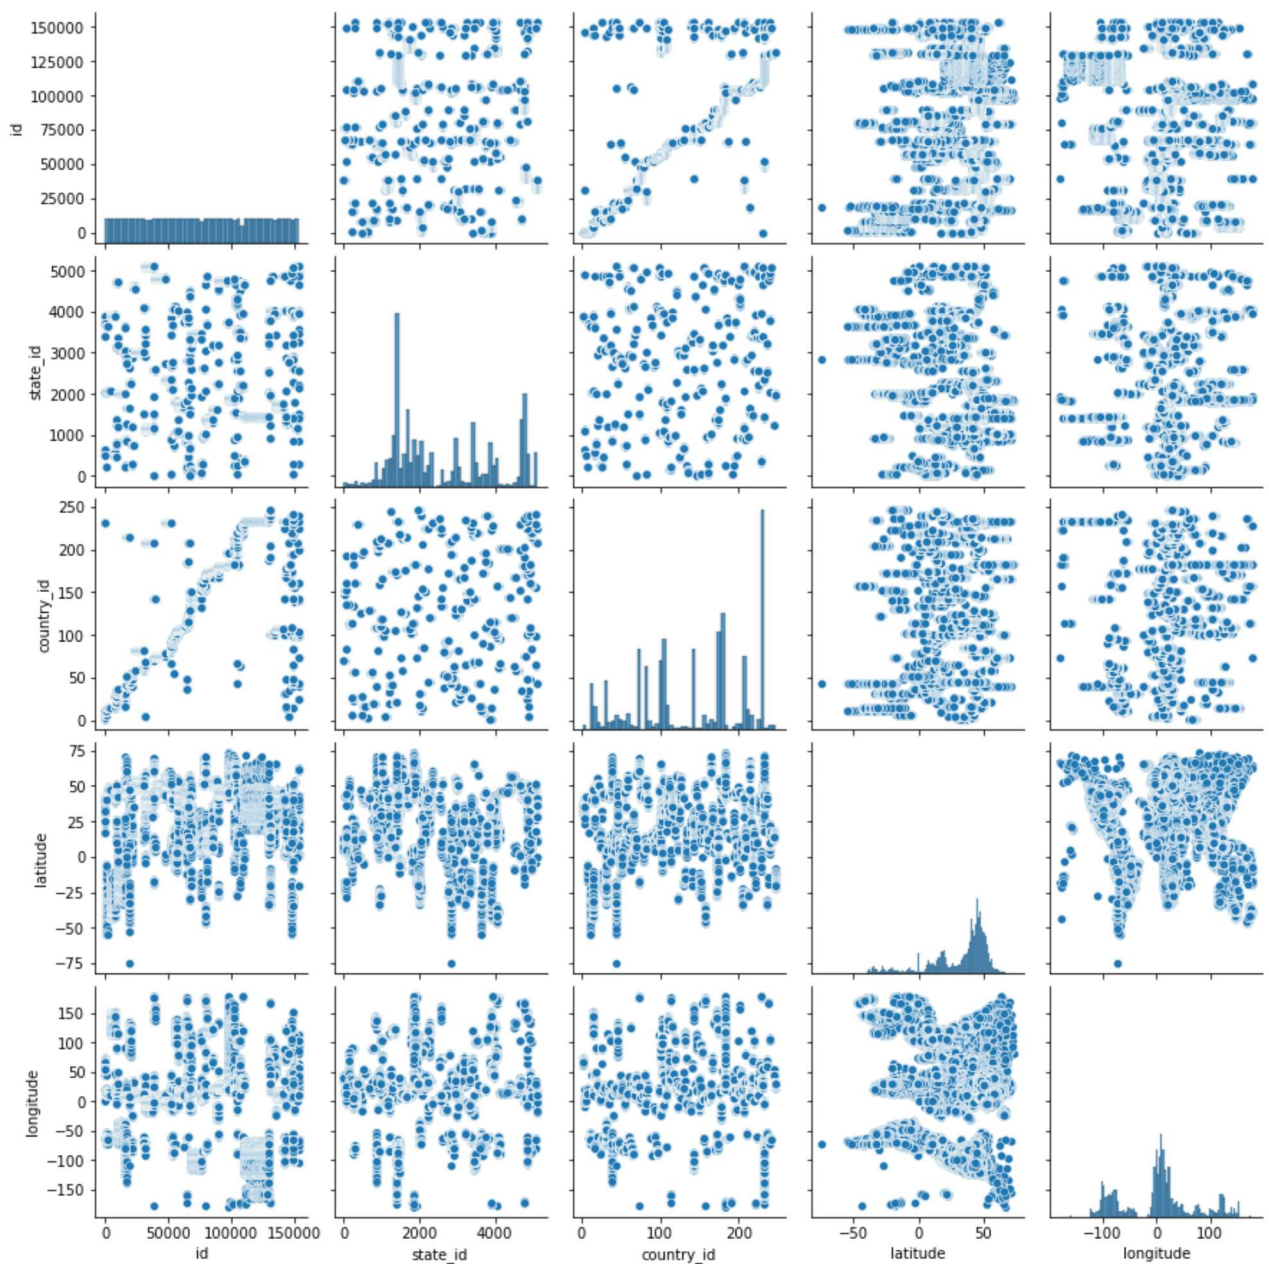
```
Out[6]: Index(['id', 'name', 'state_id', 'state_code', 'state_name', 'country_id',
              'country_code', 'country_name', 'latitude', 'longitude', 'wikiDataId'],
              dtype='object')
```

EDA and VISUALIZATION

In [7]:

```
sns.pairplot(df)
```


```
Out[7]: <seaborn.axisgrid.PairGrid at 0x239d5c4c040>
```



```
In [8]: sns.distplot(df['state_id'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

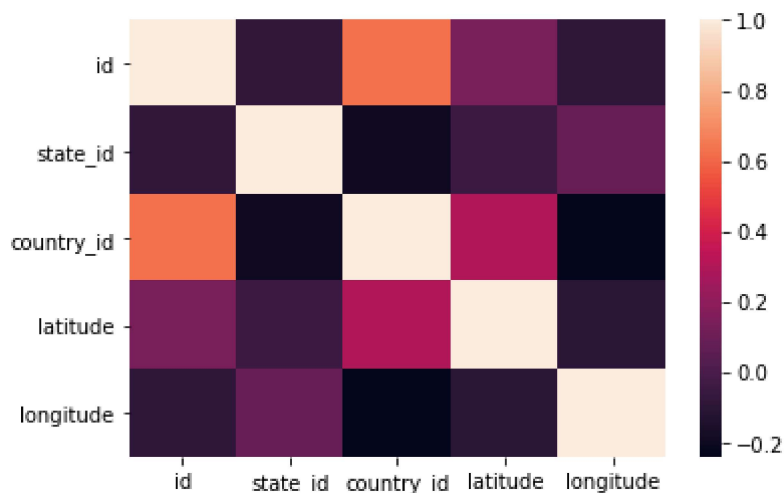
```
Out[8]: <AxesSubplot:xlabel='state_id', ylabel='Density'>
```



```
In [9]: df1 = df[['id', 'name', 'state_id', 'state_code', 'state_name', 'country_id',
                'country_code', 'country_name', 'latitude', 'longitude', 'wikiDataId']]
```

```
In [10]: sns.heatmap(df1.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: x = df1[['state_id', 'country_id',
                'latitude', 'longitude' ]]
          y = df1['id']
```

split the data into training and test data

```
In [12]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
```

```
In [13]: lr = LinearRegression()
          lr.fit(x_train, y_train)
```

```
Out[13]: LinearRegression()
```

```
In [14]: lr.intercept_
```

```
Out[14]: 16761.512014101303
```


```
In [15]: coeff = pd.DataFrame(lr.coef_, x.columns, columns =['Co-efficient'])  
coeff
```

```
Out[15]:
```

| | Co-efficient |
|-------------------|--------------|
| state_id | 1.353642 |
| country_id | 419.405326 |
| latitude | -96.578760 |
| longitude | 41.972795 |

```
In [16]: prediction = lr.predict(x_test)  
plt.scatter(y_test, prediction)
```

```
Out[16]: <matplotlib.collections.PathCollection at 0x239d78af640>
```



```
In [17]: lr.score(x_test,y_test)
```

```
Out[17]: 0.4088022617416732
```

```
In [18]: from sklearn.linear_model import Ridge,Lasso
```

```
In [19]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)  
rr.score(x_test,y_test)  
rr.score(x_train,y_train)
```

```
Out[19]: 0.4037205135331342
```

```
In [20]: rr.score(x_test,y_test)
```

```
Out[20]: 0.4088022618370707
```

```
In [21]: la = Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[21]: Lasso(alpha=10)
```

```
In [22]: la.score(x_test,y_test)
```

```
Out[22]: 0.408802357223611
```

```
In [23]: from sklearn.linear_model import ElasticNet
en = ElasticNet()
en.fit(x_train,y_train)
```

```
Out[23]: ElasticNet()
```

```
In [24]: print(en.coef_)
```

```
[ 1.3531898  419.34436461 -96.43451329  41.95872567]
```

```
In [25]: print(en.intercept_)
```

```
16766.7846160139
```

```
In [26]: print(en.predict(x_test))
```

```
[ 27226.19833714  21885.93724355  32737.32980837 ...  80897.04414038
 25583.38344341 112246.2501921 ]
```

```
In [27]: print(en.score(x_test,y_test))
```

```
0.40880275794908827
```

Evaluation Metrics

```
In [28]: from sklearn import metrics
```

```
In [29]: print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolute Error: 22616.64507184908
```

```
In [30]: print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

```
Mean Squared Error: 1156651516.1327608
```

```
In [31]: print("Root Mean Squared Error:", np.sqrt(metrics.mean_squared_error(y_test, prediction)))
```

Root Mean Squared Error: 34009.57977001128

```
In [32]: import pickle
```

```
In [33]: filename='prediction'
pickle.dump(lr, open(filename, 'wb'))
```

```
In [34]: import pandas as pd
import pickle
```

```
In [35]: filename='prediction'
model = pickle.load(open(filename, 'rb'))
```

```
In [36]: real = [[10, 20, 30, 40], [11, 45, 42, 13]]
result = model.predict(real)
```

```
In [37]: result
```

```
Out[37]: array([23944.70394288, 32138.98015648])
```