

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df = pd.read_csv("nuclear.csv").dropna()
df
```

Out[2]:

	WEAPON SOURCE COUNTRY	WEAPON DEPLOYMENT LOCATION	Data.Source	Location.Cordinates.Latitude	Location.Cordinates.Lo
0	USA	Alamogordo	DOE	32.54	
1	USA	Hiroshima	DOE	34.23	
2	USA	Nagasaki	DOE	32.45	
3	USA	Bikini	DOE	11.35	
4	USA	Bikini	DOE	11.35	
...	
2041	CHINA	Lop Nor	HFS	41.69	
2042	INDIA	Pokhran	HFS	27.07	
2043	INDIA	Pokhran	NRD	27.07	
2044	PAKIST	Chagai	HFS	28.90	
2045	PAKIST	Kharan	HFS	28.49	

2046 rows × 6 columns

```
In [3]: df.head()
```

Out[3]:

	WEAPON SOURCE COUNTRY	WEAPON DEPLOYMENT LOCATION	Data.Source	Location.Cordinates.Latitude	Location.Cordinates.Longit
0	USA	Alamogordo	DOE	32.54	-10
1	USA	Hiroshima	DOE	34.23	13
2	USA	Nagasaki	DOE	32.45	12
3	USA	Bikini	DOE	11.35	16
4	USA	Bikini	DOE	11.35	16

Data cleaning and pre processing

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2046 entries, 0 to 2045
Data columns (total 16 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   WEAPON SOURCE COUNTRY                 2046 non-null   object
 1   WEAPON DEPLOYMENT LOCATION           2046 non-null   object
 2   Data.Source                           2046 non-null   object
 3   Location.Cordinates.Latitude         2046 non-null   float64
 4   Location.Cordinates.Longitude        2046 non-null   float64
 5   Data.Magnitude.Body                  2046 non-null   float64
 6   Data.Magnitude.Surface                2046 non-null   float64
 7   Location.Cordinates.Depth             2046 non-null   float64
 8   Data.Yeild.Lower                      2046 non-null   float64
 9   Data.Yeild.Upper                      2046 non-null   float64
10   Data.Purpose                            2046 non-null   object
11   Data.Name                             2046 non-null   object
12   Data.Type                             2046 non-null   object
13   Date.Day                              2046 non-null   int64
14   Date.Month                            2046 non-null   int64
15   Date.Year                             2046 non-null   int64
dtypes: float64(7), int64(3), object(6)
memory usage: 271.7+ KB
```

In [5]: `df.describe()`

Out[5]:

	Location.Cordinates.Latitude	Location.Cordinates.Longitude	Data.Magnitude.Body	Data.Mag
count	2046.000000	2046.000000	2046.000000	
mean	35.462429	-36.015037	2.145406	
std	23.352702	100.829355	2.625453	
min	-49.500000	-169.320000	0.000000	
25%	37.000000	-116.051500	0.000000	
50%	37.100000	-116.000000	0.000000	
75%	49.870000	78.000000	5.100000	
max	75.100000	179.220000	7.400000	

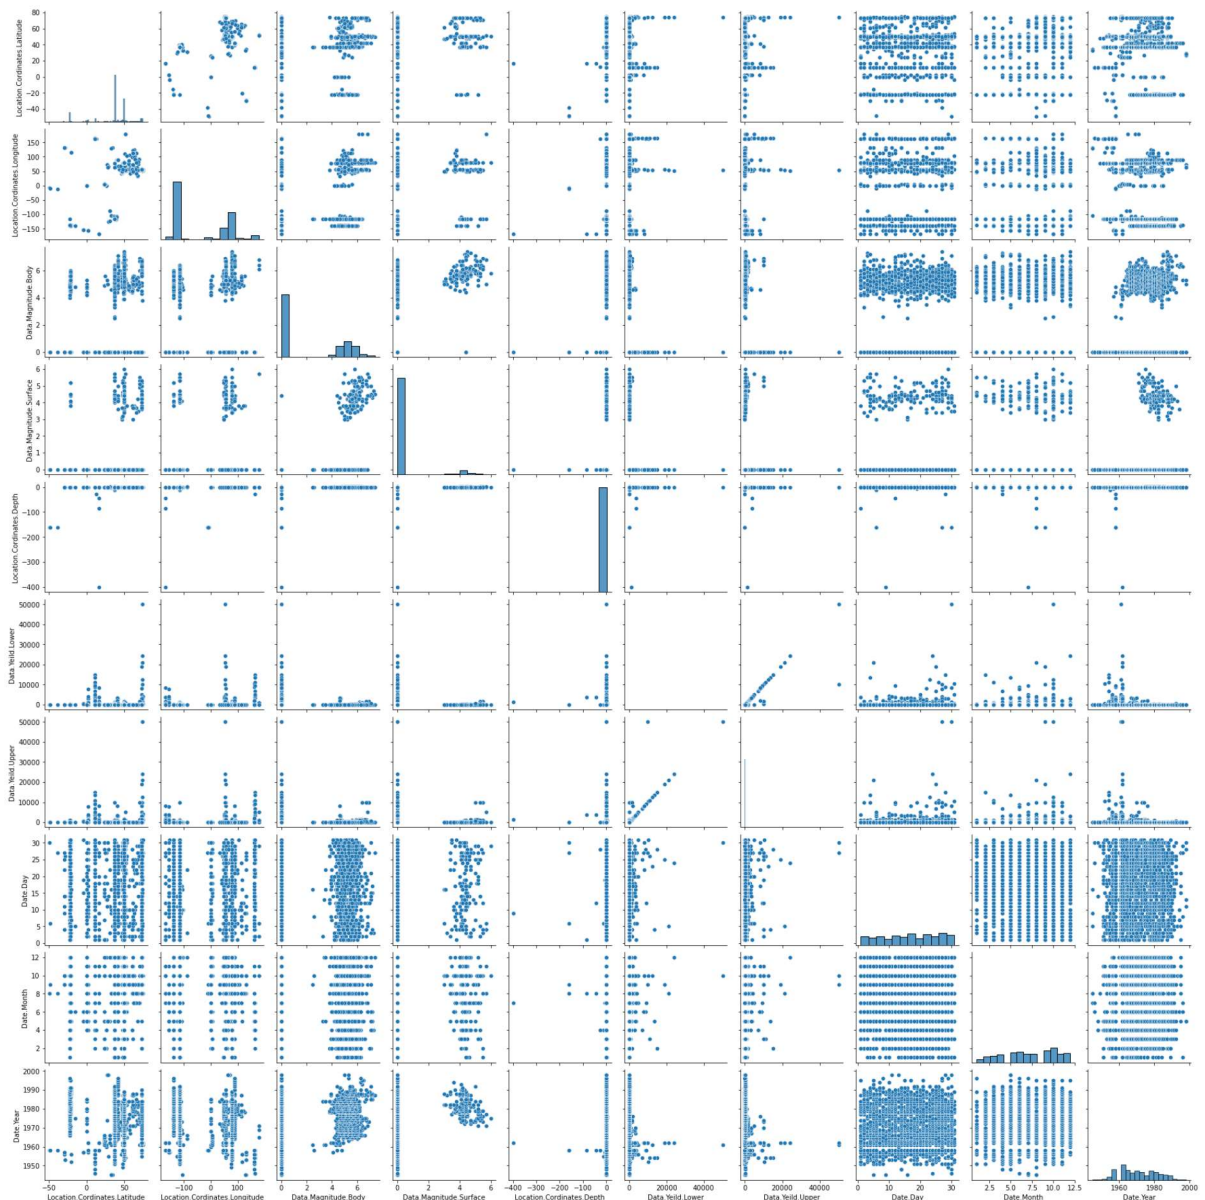
```
In [6]: df.columns
```

```
Out[6]: Index(['WEAPON_SOURCE_COUNTRY', 'WEAPON_DEPLOYMENT_LOCATION', 'Data.Source',  
              'Location.Coordinates.Latitude', 'Location.Coordinates.Longitude',  
              'Data.Magnitude.Body', 'Data.Magnitude.Surface',  
              'Location.Coordinates.Depth', 'Data.Yeild.Lower', 'Data.Yeild.Upper',  
              'Data.Purpose', 'Data.Name', 'Data.Type', 'Date.Day', 'Date.Month',  
              'Date.Year'],  
             dtype='object')
```

EDA and VISUALIZATION

```
In [7]: sns.pairplot(df)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x173014c1820>
```

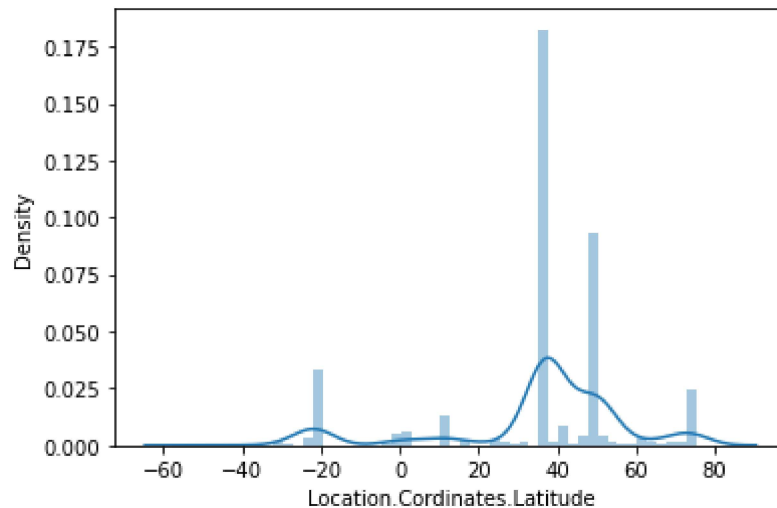


```
In [8]: sns.distplot(df['Location.Cordinates.Latitude'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

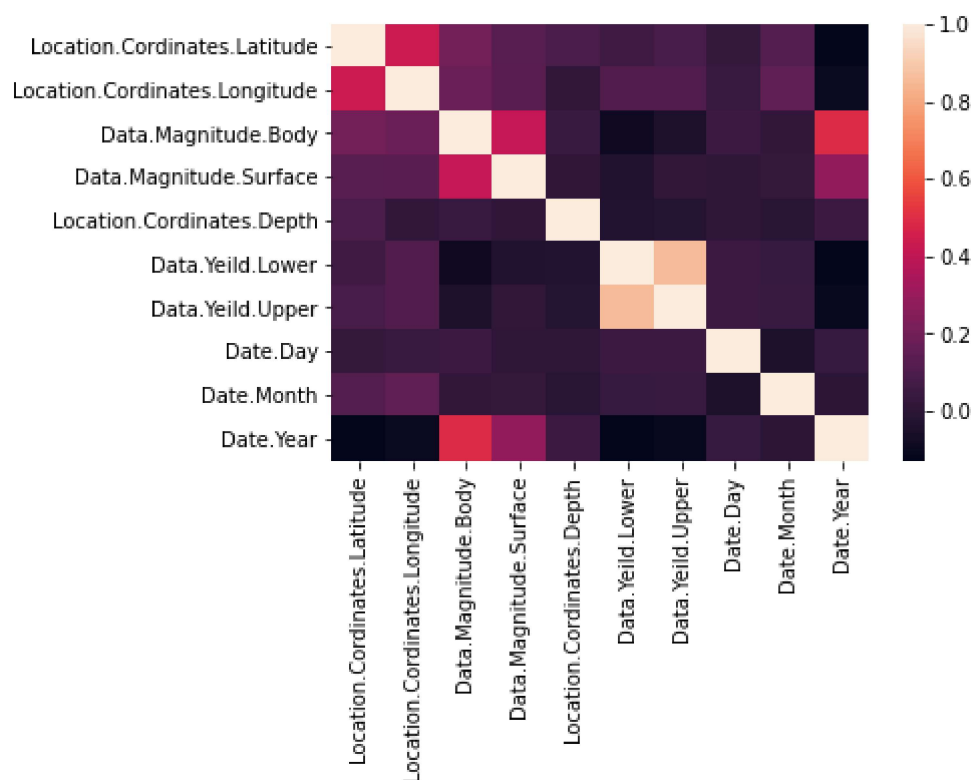
```
Out[8]: <AxesSubplot:xlabel='Location.Cordinates.Latitude', ylabel='Density'>
```



```
In [9]: df1 = df[['WEAPON SOURCE COUNTRY', 'WEAPON DEPLOYMENT LOCATION', 'Data.Source',  
                'Location.Cordinates.Latitude', 'Location.Cordinates.Longitude',  
                'Data.Magnitude.Body', 'Data.Magnitude.Surface',  
                'Location.Cordinates.Depth', 'Data.Yeild.Lower', 'Data.Yeild.Upper',  
                'Data.Purpose', 'Data.Name', 'Data.Type', 'Date.Day', 'Date.Month',  
                'Date.Year']]
```

```
In [10]: sns.heatmap(df1.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: x = df1[['Location.Cordinates.Latitude', 'Location.Cordinates.Longitude',
                  'Data.Magnitude.Body',
                  'Location.Cordinates.Depth', 'Data.Yeild.Lower', 'Data.Yeild.Upper', 'Date.Day',
                  'Date.Month', 'Date.Year']]
y = df1['Data.Magnitude.Surface']
```

split the data into training and test data

```
In [12]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
```

```
In [13]: lr = LinearRegression()
lr.fit(x_train, y_train)
```

```
Out[13]: LinearRegression()
```

```
In [14]: lr.intercept_
```

```
Out[14]: -34.10720923778148
```


```
In [15]: coeff = pd.DataFrame(lr.coef_, x.columns, columns =['Co-efficient'])
coeff
```

Out[15]:

	Co-efficient
Location.Cordinates.Latitude	0.003430
Location.Cordinates.Longitude	0.000812
Data.Magnitude.Body	0.147027
Location.Cordinates.Depth	-0.001198
Data.Yeild.Lower	-0.000049
Data.Yeild.Upper	0.000045
Date.Day	-0.003632
Date.Month	0.003313
Date.Year	0.017296

```
In [16]: prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x17312755670>



```
In [17]: lr.score(x_test,y_test)
```

Out[17]: 0.19921989735087386

```
In [18]: from sklearn.linear_model import Ridge,Lasso
```

```
In [19]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
rr.score(x_train,y_train)
```

Out[19]: 0.1948024994212233

In [20]: `rr.score(x_test,y_test)`

Out[20]: 0.19921997008820713

In [21]: `la = Lasso(alpha=10)`
`la.fit(x_train,y_train)`

Out[21]: Lasso(alpha=10)

In [22]: `la.score(x_test,y_test)`

Out[22]: 0.017800264797578902

In [23]: `from sklearn.linear_model import ElasticNet`
`en = ElasticNet()`
`en.fit(x_train,y_train)`

Out[23]: ElasticNet()

In [24]: `print(en.coef_)`

```
[ 4.13404901e-03  1.25583655e-03  5.13577239e-02 -0.00000000e+00
 -7.68123337e-05  5.26147310e-05 -0.00000000e+00  0.00000000e+00
  2.42579176e-02]
```

In [25]: `print(en.intercept_)`

```
-47.667432481886046
```

In [26]: `print(en.predict(x_test))`

```
[ 2.26031951e-01 -1.63147026e-01 -6.50630608e-02  8.33484951e-02
  7.56102611e-01  9.76700058e-01  6.44255909e-01 -3.33702575e-01
  5.76643363e-01  3.31232264e-01  1.55859706e-02  7.32735069e-01
  2.82889056e-01  4.71869158e-01  3.10424715e-02  8.47925579e-02
  2.31259625e-01  5.52484350e-01  2.22126915e-03  5.71104038e-02
  2.01774033e-01  2.31119226e-01 -4.08051431e-02  4.34963737e-01
 -3.83084374e-01 -3.56183068e-01  4.71728759e-01 -1.54218797e-01
  1.53258198e-01  9.23463959e-01 -1.63679373e-01  2.87860923e-01
  2.31201498e-01  8.36295005e-01  8.80134862e-01  3.70513710e-01
  7.03211295e-01 -3.32976647e-01  1.31411030e+00 -1.65472255e-02
  8.86071222e-01  7.56382277e-01  2.31258822e-01  2.01774033e-01
  7.99183670e-02  7.11737239e-01 -1.65472255e-02  6.19386307e-01
  5.71922637e-01  3.17627163e-01  7.26649733e-01  3.17846127e-01
  1.04357740e-01  1.95829244e-01  7.27195821e-01  4.80125474e-03
  8.98958349e-01  1.29273981e-01  5.50914268e-01  9.13277378e-01
  1.05468052e+00 -6.76360184e-02  2.48418856e-01 -2.67377866e-01
  6.68436373e-01  6.63762835e-01  8.04970033e-02  3.10714378e-01
  3.55574944e-01  9.65432603e-04  1.75394548e-01  6.43644224e-01
 -1.63146973e-01  5.51342695e-01  8.23830107e-02  1.75545424e-01
  4.06500404e-01  3.00000000e-01  6.05445070e-01  5.00000000e-01]
```

```
In [27]: print(en.score(x_test,y_test))
```

0.17049406052259697

Evaluation Metrics

```
In [28]: from sklearn import metrics
```

```
In [29]: print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolute Error: 0.6363042556370765

```
In [30]: print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

Mean Squared Error: 1.118977148440899

```
In [31]: print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root Mean Squared Error: 1.057817162103593

```
In [32]: import pickle
```

```
In [33]: filename='prediction'
pickle.dump(lr,open(filename,'wb'))
```

```
In [34]: import pandas as pd
import pickle
```

```
In [35]: filename='prediction'
model = pickle.load(open(filename,'rb'))
```

```
In [37]: real = [[10,20,30,44,55,66,99,77,88],[11,45,55,12,23,56,45,89,78]]
result = model.predict(real)
```

```
In [38]: result
```

```
Out[38]: array([-28.28074631, -24.47893551])
```