

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [2]:

```
df = pd.read_csv("states.csv").dropna()
df
```

Out[2]:

	id	name	country_id	country_code	country_name	state_code	type	latitude	long
170	3656	Buenos Aires	11	AR	Argentina	B	province	-37.201729	-59.841
171	3647	Catamarca	11	AR	Argentina	K	province	-28.471588	-65.787
172	3640	Chaco	11	AR	Argentina	H	province	-27.425718	-59.024
173	3651	Chubut	11	AR	Argentina	U	province	-43.293425	-65.111
Ciudad Autónoma de Buenos Aires			11	AR	Argentina	C	city	-34.603684	-58.386
...
4968	2041	Yaracuy	239	VE	Venezuela	U	state	10.339389	-68.814
4969	2042	Zulia	239	VE	Venezuela	V	state	10.291024	-72.146
5033	5074	Saint Croix	242	VI	Virgin Islands (US)	SC	district	17.729352	-64.735
5034	5073	Saint John	242	VI	Virgin Islands (US)	SJ	district	18.335617	-64.806
5035	5072	Saint Thomas	242	VI	Virgin Islands (US)	ST	district	18.342846	-65.071

1580 rows × 9 columns



In [3]:

```
df.head()
```

Out[3]:

	id	name	country_id	country_code	country_name	state_code	type	latitude	long
170	3656	Buenos Aires	11	AR	Argentina	B	province	-37.201729	-59.841
171	3647	Catamarca	11	AR	Argentina	K	province	-28.471588	-65.787
172	3640	Chaco	11	AR	Argentina	H	province	-27.425718	-59.024
173	3651	Chubut	11	AR	Argentina	U	province	-43.293425	-65.111

	id	name	country_id	country_code	country_name	state_code	type	latitude	longit
174	4880	Ciudad Autónoma de Buenos Aires	11	AR	Argentina	C	city	-34.603684	-58.381

Data cleaning and pre processing

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1580 entries, 170 to 5035
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   id          1580 non-null    int64  
 1   name         1580 non-null    object  
 2   country_id   1580 non-null    int64  
 3   country_code 1580 non-null    object  
 4   country_name 1580 non-null    object  
 5   state_code   1580 non-null    object  
 6   type         1580 non-null    object  
 7   latitude     1580 non-null    float64 
 8   longitude    1580 non-null    float64 
dtypes: float64(2), int64(2), object(5)
memory usage: 123.4+ KB
```

In [5]: `df.describe()`

Out[5]:

	id	country_id	latitude	longitude
count	1580.000000	1580.000000	1580.000000	1580.000000
mean	2685.916456	134.000633	26.988930	15.009671
std	1611.169440	66.055166	19.635279	66.200355
min	48.000000	11.000000	-54.805400	-178.116500
25%	1339.750000	75.000000	13.752013	-7.622388
50%	2210.500000	139.000000	30.887089	11.665277
75%	4013.250000	178.000000	42.938004	45.682217
max	5220.000000	242.000000	77.874972	166.649935

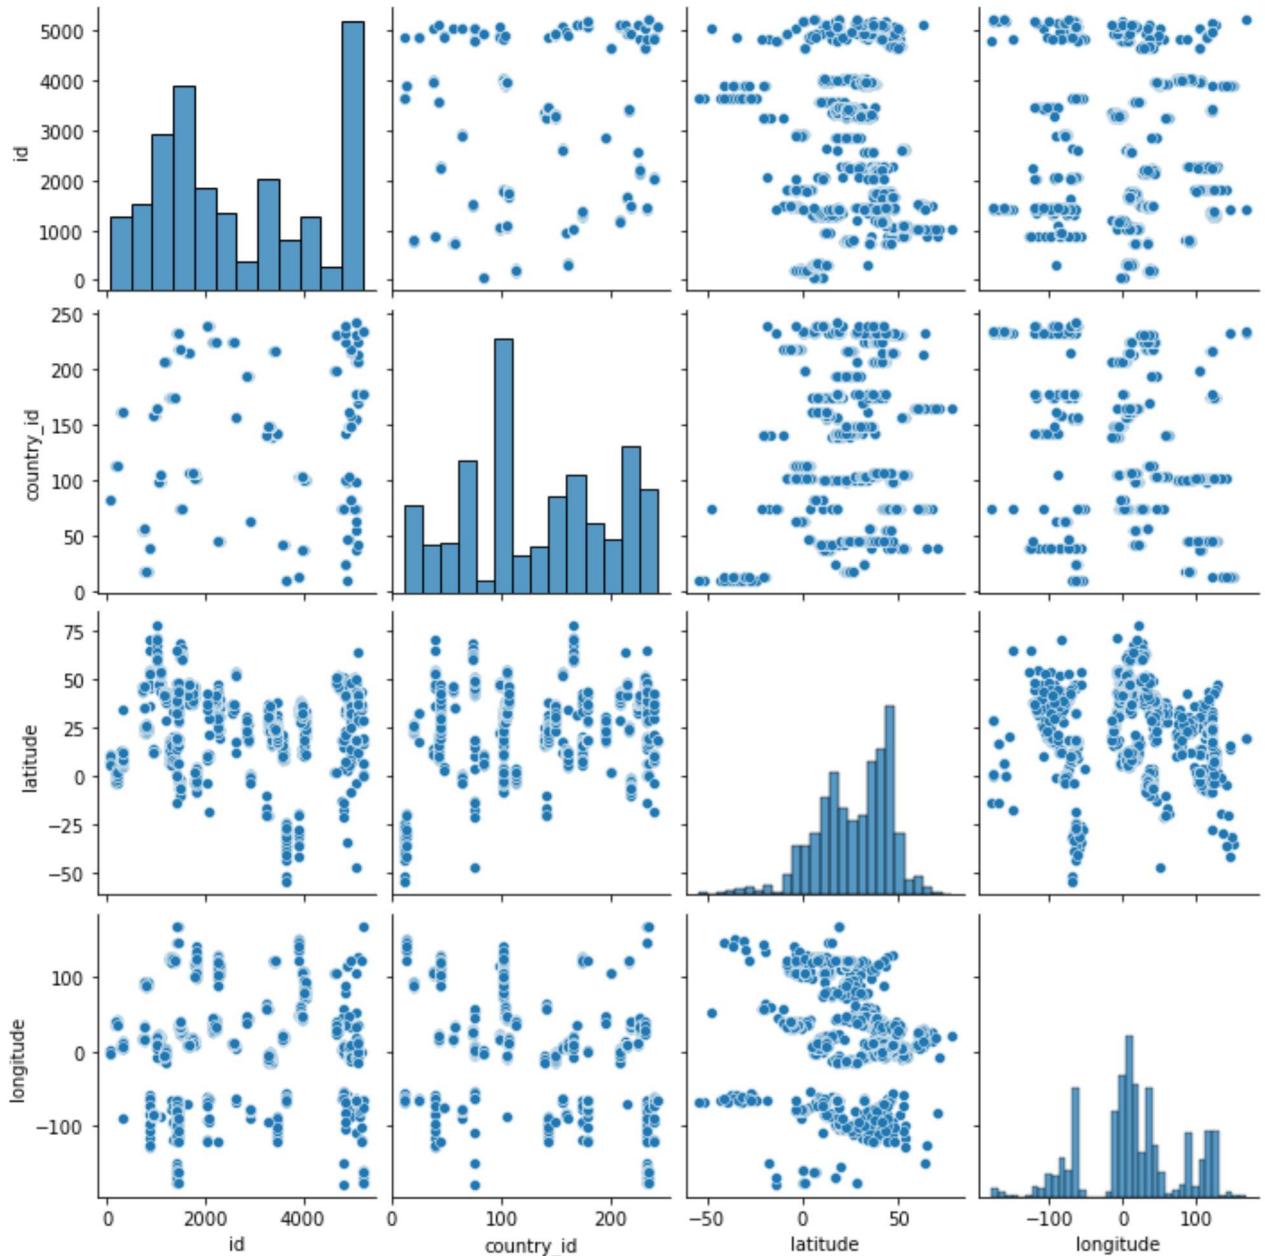
In [6]: `df.columns`

Out[6]: `Index(['id', 'name', 'country_id', 'country_code', 'country_name', 'state_code', 'type', 'latitude', 'longitude'], dtype='object')`

EDA and VISUALIZATION

In [7]:
`sns.pairplot(df)`

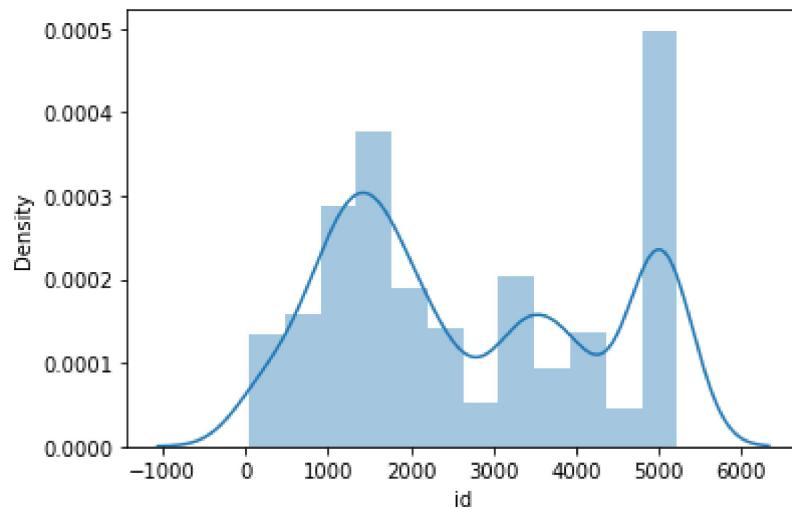
Out[7]: <seaborn.axisgrid.PairGrid at 0x1a0dfa784c0>



In [8]:
`sns.distplot(df['id'])`

```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning:
`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)
```

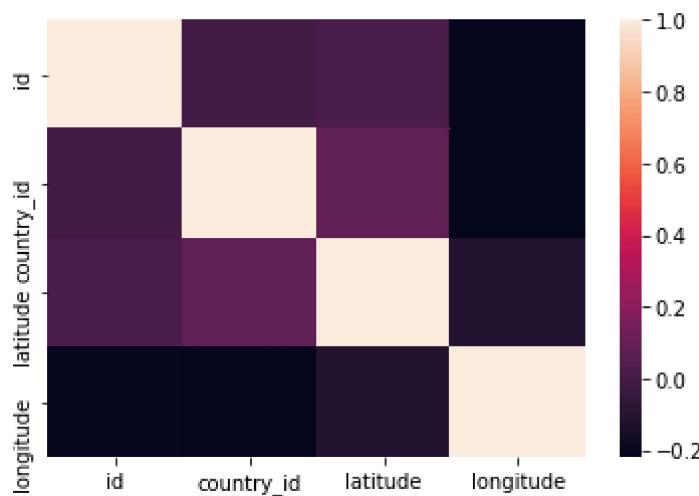
Out[8]: <AxesSubplot:xlabel='id', ylabel='Density'>



```
In [9]: df1 = df[['id', 'name', 'country_id', 'country_code', 'country_name',
   'state_code', 'type', 'latitude', 'longitude']]
```

```
In [10]: sns.heatmap(df1.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: x = df1[['id', 'country_id', 'latitude', ]]
y = df1['longitude']
```

split the data into training and test data

```
In [12]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
```

```
In [13]: lr = LinearRegression()
lr.fit(x_train, y_train)
```

```
Out[13]: LinearRegression()
```

```
In [14]: lr.intercept_
```

```
Out[14]: 71.72030586537088
```

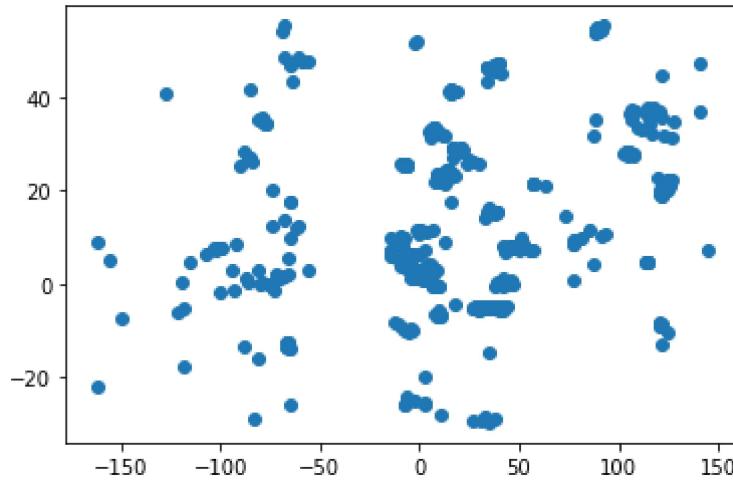
```
In [15]: coeff = pd.DataFrame(lr.coef_, x.columns, columns =['Co-efficient'])
coeff
```

```
Out[15]: Co-efficient
```

	Co-efficient
id	-0.008164
country_id	-0.211728
latitude	-0.285410

```
In [16]: prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

```
Out[16]: <matplotlib.collections.PathCollection at 0x1a0e0e7e310>
```



```
In [17]: lr.score(x_test,y_test)
```

```
Out[17]: 0.09089056655038252
```

```
In [18]: from sklearn.linear_model import Ridge,Lasso
```

```
In [19]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
rr.score(x_train,y_train)
```

```
Out[19]: 0.09309894107374084
```

```
In [20]: rr.score(x_test,y_test)
```

```
Out[20]: 0.09089058804969075
```

```
In [21]: la = Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[21]: Lasso(alpha=10)
```

```
In [22]: la.score(x_test,y_test)
```

```
Out[22]: 0.09090171311470485
```

```
In [23]: from sklearn.linear_model import ElasticNet
en = ElasticNet()
en.fit(x_train,y_train)
```

```
Out[23]: ElasticNet()
```

```
In [24]: print(en.coef_)
```

```
[-0.00816375 -0.21161806 -0.28376439]
```

```
In [25]: print(en.intercept_)
```

```
71.66111349722851
```

```
In [26]: print(en.predict(x_test))
```

```
[ 2.78066968 -13.21948028  36.46321718  41.15978801  26.52965804
 -0.18612289   7.45059281   7.69820457   0.39386783  -4.61684436
 32.63781437   8.64297654   6.60212562   4.24060313 -13.48450436
 -8.58106823  36.81994093  55.30634239  22.12197696   9.60294091
 54.40705442 -13.56469124  47.39971155  54.33481076 -26.00427662
 7.74436764   6.44351753 -13.71676975  25.9110113  23.662018
 4.637796   22.82705143  28.21326034  51.51259539  -5.33283945
 -8.83103889  26.50344289  54.77570774  14.54685719  54.52034508
 2.80298492   2.16878783  -5.36635404 -25.54103606   0.35787668
 -4.88111676  -4.91882823  26.32297525  7.56323084  8.65859792
 -14.85672489  47.32397832  23.38332928 -25.19609591 -17.90940696
 6.14250815  21.22129581  35.52288211  46.41635301  8.86205033
 -0.56272691 -13.27768735 -0.10848745 -0.31594314  3.04218064
 27.00935124   1.32673972  1.26872908  22.30001376 -4.60079487
 54.75500787  21.10768738 -13.63530044  21.70270835 -6.7953903
 -7.29305524  32.38932881  4.45244629 -25.71596151  35.03003631
 20.06089024  53.9263728   4.70502564  15.40909707 11.15021547
 -4.87144371  25.54916743  34.92736638  25.54093938 47.42838155
 8.26791806  29.15494426  21.87227759 -0.69082123 21.81436722
 -9.36153451   9.69667814  -6.63948466  22.62969249 17.55965739
 24.3554651  -0.45246768  22.96102456  46.33286193 41.16230048
 36.10379411   1.60823016  22.26954694  1.7183622  54.4965484
 -29.25003069   8.35114393  48.613825  -8.32391381 31.7118988
 -6.00075605   1.61261288  -8.70163573  1.38800198 40.9888288
 -4.9544191  20.31897181  17.4688299  29.11956081 -13.24453026
 31.59675993  19.01729861  21.77243204  35.68098535 -12.72210212
 -24.34692436  -0.26452312  46.37103124 -13.25959224 19.21384818
 2.34373249  21.86180724  2.73654979  4.63948628  2.15145741
```

2.19384829	0.96482001	4.59298341	4.67187928	1.90689627
1.55932317	54.18272061	12.24266777	25.36124097	26.83544633
-4.86572646	2.78724267	-4.87300043	20.18533885	14.89424651
46.2473332	54.00313488	7.05001026	47.25292124	22.55210557
4.67963249	22.20823744	-4.30267297	20.11454819	23.42391973
53.92561122	20.44918262	3.46964187	-22.25098826	34.76453015
28.78248667	20.62048831	-6.35601805	54.70126213	3.45790471
18.8379576	28.15919971	27.42153542	-6.73200332	34.36847069
-1.16638631	-0.5953149	21.25503454	43.47141039	41.1530576
-4.9387612	2.11941378	21.0137115	45.7134753	28.36647911
27.81665963	22.18687987	41.55974551	22.10537278	33.2537554
54.08696546	3.64896559	-9.78862623	47.18867553	-19.72709773
21.82997646	36.80903909	47.25243535	41.79420635	-0.16310112
-9.78837385	-5.19056675	-5.53955132	41.01901354	7.52102661
-0.56097004	17.50833614	4.02088329	31.72080303	9.67555576
25.54794411	48.57194906	25.84928098	7.53725178	28.3660859
-27.85098018	7.10331515	21.1417075	-5.14987601	28.8704378
-13.03048916	8.21182679	6.31163147	-8.66155723	-28.48350996
37.00445273	47.24400565	46.21174467	11.4217743	25.9464007
21.56516888	52.11159554	-13.42124585	2.0667634	1.74963176
4.12650369	23.01283744	-28.7172263	36.03442802	2.62122002
10.15310883	7.17490863	54.58556054	9.84409703	54.84286929
-13.66729878	10.62036522	45.20753494	54.80286673	2.90662268
46.02787427	-15.99030119	-0.50289552	-5.75096666	4.61536928
5.22894514	28.66323946	3.93454813	2.95550066	1.29351714
28.31946373	24.58497263	9.3160659	0.89953924	6.58350077
20.38564429	34.35940447	-13.3294921	5.98033009	22.24015015
-5.49638538	34.97658257	6.6896295	7.80570132	8.26418679
28.03826162	-1.91949381	0.28354207	46.57194363	45.47137129
43.59373592	-5.08839618	4.03856347	-5.1548879	29.11884829
27.90300634	53.93939663	4.64944733	-12.66806306	54.54229897
11.75158853	9.16741175	-10.32318209	31.50272563	25.43512461
21.08636782	9.83495206	25.27119474	-0.5810521	21.20992565
7.01187787	8.00509822	37.72242572	7.20941969	21.37570838
41.59733922	2.55385095	-0.41656564	-0.21006161	23.13147286
10.55237231	12.22859958	20.16573678	22.89362771	54.19605026
13.76455768	22.52459034	33.1690973	-5.4650817	-7.08764867
0.89150675	37.5358143	-29.7039955	20.49490767	8.0524344
21.91826806	22.29172237	-4.84581016	-8.19195043	28.50991124
54.97013435	54.77271607	7.06772111	-4.66528363	5.18906258
-5.4088676	6.08180676	-0.26029615	4.57439132	11.58122911
8.01618493	1.55417664	7.18717705	-4.71533342	34.31335948
-25.90266889	40.99660489	3.7396178	28.67492453	21.83586581
15.25464665	10.35707553	0.5941295	1.49379648	14.1542731
27.21554305	21.04502417	27.32557713	35.09912361	35.19688708
26.58601022	54.38470858	26.41044386	46.20292694	33.54670933
27.66994376	-13.41027613	37.11742915	46.37829351	8.39520893
-4.73453825	21.19632848	-9.45755189	-12.64455158	5.06596668
-4.58457268	1.38129834	-13.14172883	36.58508665	24.29148874
2.41737398	-9.86444672	1.29692433	11.3547184	47.36739934
22.69820483	3.127976	7.69989778	-10.31743039	6.79414583
15.29181936	46.31057098	3.96747025	27.67509997	54.841145
25.41107269	28.72027525	55.08745304	53.65463372	44.66341819
19.45006943	3.03391836	-1.44028957	32.67480757	-8.98519767
31.3057446	33.27391809	-25.71958894	-13.51191256	54.36297852
-4.98031357	19.83026044	-29.34715558	5.97168142	24.08196148
2.36226849	31.98565366	33.00784936	2.70074329	-4.99923298
32.50935439	55.17509876	16.06865384	54.63874985	-5.37598594
19.06681327	4.15673684	19.11929655	8.23087177	22.15972434
1.23290671	4.14619377	19.16660415	-5.2879037	-5.00191482
29.00897788	-13.3527321	0.72113976	9.70128327	2.18239976
-0.51382342	25.60987695	4.62612687	23.27508505	-9.76256798
-0.48415983	3.20153835	31.67740843	0.40541578	5.3672734
46.55563919	11.79757531	21.8886819	4.14427648	25.9009335
6.48648088	46.63006956	-5.19626551	22.02666406	33.56969806

```
1.29838001 -28.9323833 -5.45128822 25.50835792 45.95446249
34.26383232 2.90734448 25.62646606 -6.99453228]
```

In [27]:

```
print(en.score(x_test,y_test))
```

0.09089540832542065

Evaluation Metrics

In [28]:

```
from sklearn import metrics
```

In [29]:

```
print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolute Error: 47.73228923464871

In [30]:

```
print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

Mean Squared Error: 3835.1091966366057

In [31]:

```
print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root Mean Squared Error: 61.9282584660396

In [32]:

```
import pickle
```

In [33]:

```
filename='prediction'
pickle.dump(lr,open(filename,'wb'))
```

In [34]:

```
import pandas as pd
import pickle
```

In [35]:

```
filename='prediction'
model = pickle.load(open(filename,'rb'))
```

In [37]:

```
real = [[10,20,30],[11,45,55]]
result = model.predict(real)
```

In [38]:

```
result
```

Out[38]:

```
array([58.84180743, 46.40519511])
```